

Bigstream Alignment Library for the RNR-ExM Challenge

ISBI 2023

Greg M. Fleishman

Scientific Software Engineer, Janelia Research Campus HHMI

Janelia Research Campus - HHMI

- Physical campus of the Howard Hughes Medical Institute
- Focus on Neuroscience and Cell biology
- Focus on new technology development
 - Microscopy, molecular tools, **computational methods**
- Scientific Computing Department
 - Core facility providing computational expertise to labs



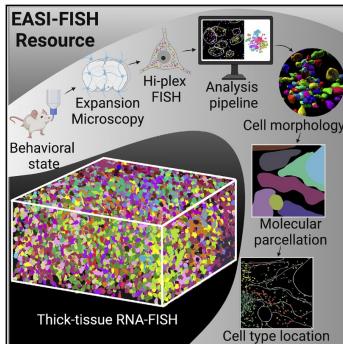
The screenshot shows the homepage of the Janelia Research Campus website. At the top, there is a navigation bar with links for NEWS, CAREERS, CONFERENCES, PEOPLE, PUBLICATIONS, and RESEARCH GROUPS. Below the navigation bar is a large banner image of a modern glass-walled building with a view of green trees and a path outside. Overlaid on the banner is the text "Small labs, big science." In the center of the page, there are two main sections: "Publications" and "Featured". The "Publications" section includes links to the PODGORSKI LAB, SCHREITER LAB, and GENIE, as well as CELL AND TISSUE CULTURE, MOLECULAR BIOLOGY, and VIRAL TOOLS. It also features a news item about a general approach to engineer positive-going eFRET voltage indicators. The "Featured" section includes a graphic showing a sensor with a color scale from 0.94 to 2.28, a news item about the newest Voltron sensor powers up its detection of smaller brain cell signals, and a 3D ribbon model of a protein structure.

What is Bigstream?

Cell

EASI-FISH for thick tissue defines lateral hypothalamus spatio-molecular organization

Graphical abstract



Resource



- Library of tools for 3D registration
- Extremely customizable pipelines
- Scales to clusters and images too-large-for-memory

- Used in several publications already
- Currently being used in projects at:
 - Janelia, UCLA, MIT, CIBR, E11 Bio, Harvard, Allen institute
- Used on data from *C. elegans*, Mouse, and Zebrafish

- **Real contribution:** Combining SimpleITK and Dask to scale all SITK registrations; and the logic to stitch things back together seamlessly.

bigstream Public

master 2 branches 2 tags Go to file Add file Code About

fleishman.github.snafoo ... 11e1338 2 days ago 102 commits

bigstream beginnings of ome-ngff for stitch results 2 days ago

notebooks updated easifish pipeline tutorial notebook 6 months ago

resources downsample certificate 2 months ago

.gitignore no more hanging clusters 2 years ago

LICENSE.txt initial commit 3 years ago

README.md add rnrexm certificate to readme 2 months ago

setup.py added h5py dependency last month

README.md

BigStream

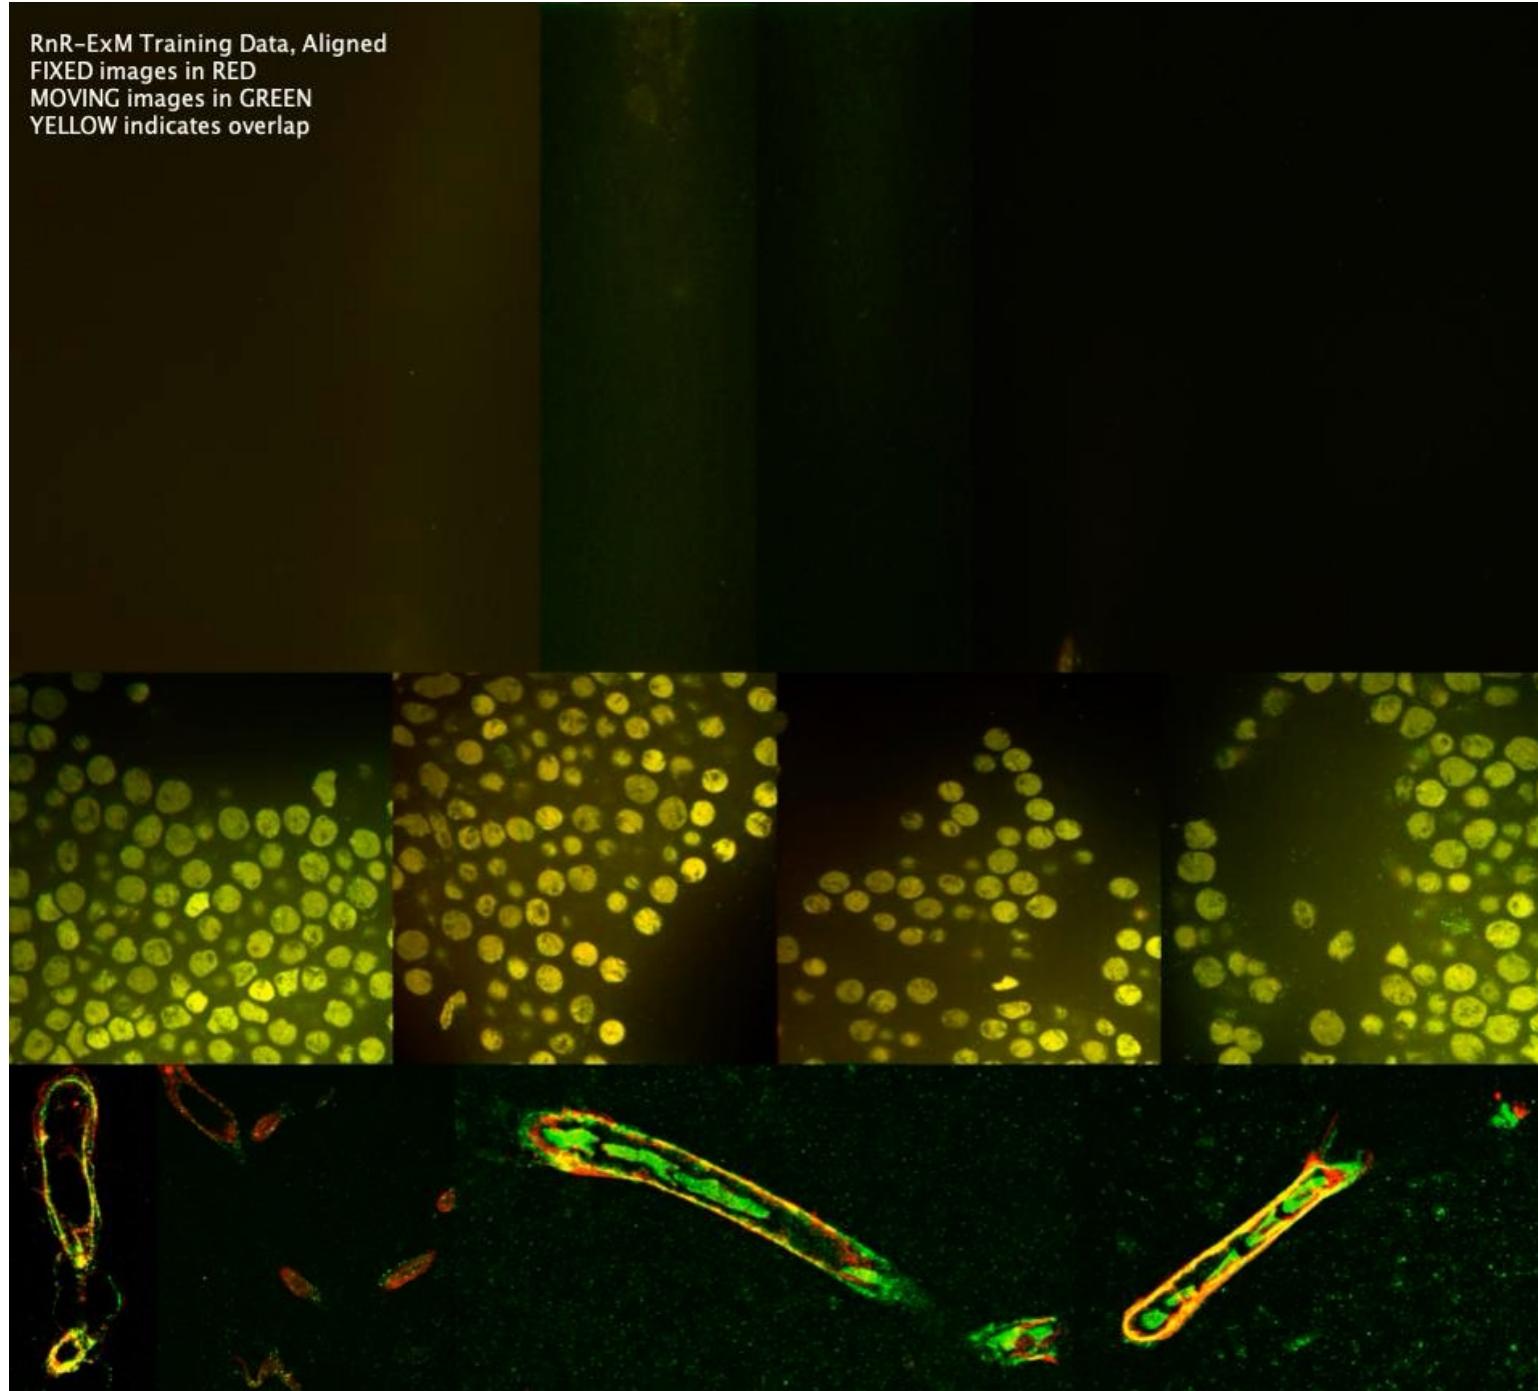


Training Data: aligned

C. elegans

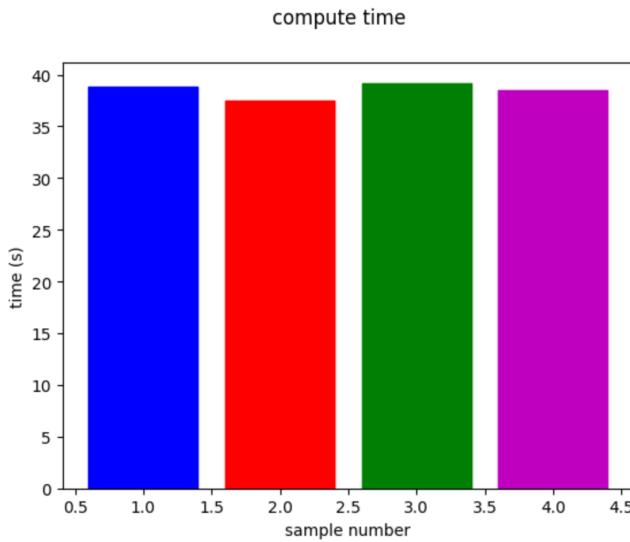
Zebrafish

Mouse

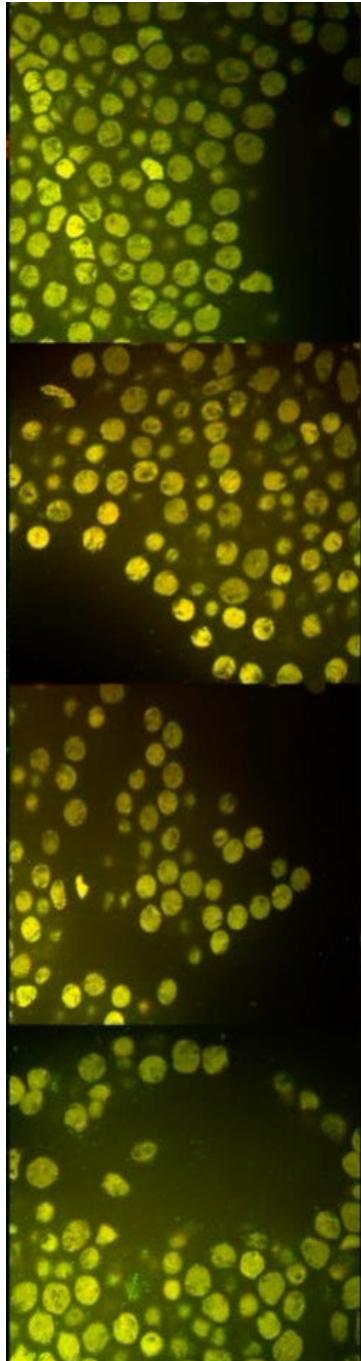




Training Data: Zebrafish



16 cores: 2.7 GHz Intel SkyLake Platinum 8168
Oracle Linux 8.3
0 GPUs



```
%%time

# alignment functions
from bigstream.align import alignment_pipeline
from bigstream.transform import apply_transform

# define alignment steps
common_kwargs = {
    'alignment_spacing':4.0,
    'shrink_factors':(2,),
    'smooth_sigmas':(2.,),
    'optimizer_args':{
        'learningRate':0.25,
        'minStep':0.,
        'numberOfIterations':400,
    },
}

affine_kwargs = {
    'initial_condition':'CENTER',
}

steps = [('affine', {**common_kwargs, **affine_kwargs}),]

# align
affine = alignment_pipeline(
    fix, mov,
    fix_spacing, mov_spacing,
    steps,
)

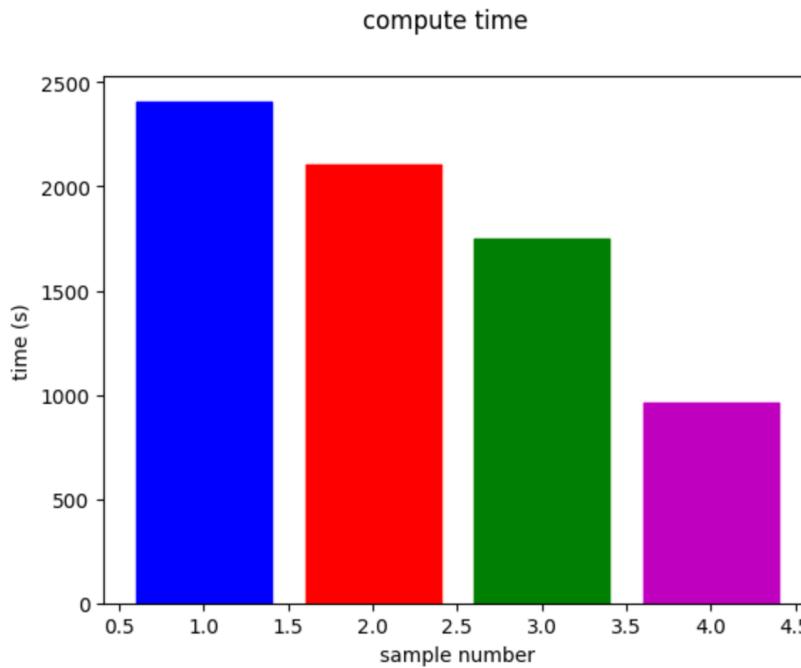
# apply affine only
affine_aligned = apply_transform(
    fix, mov,
    fix_spacing, mov_spacing,
    transform_list=[affine],
)

# write results
np.savetxt('affine.mat', affine)
nrrd.write('./affine.nrrd', affine_aligned.transpose(2,1,0), compression_level=2)

# load precomputed results
affine = np.loadtxt('./affine.mat')
```



Training Data: Mouse



16 cores: 2.7 GHz Intel SkyLake Platinum 8168
Oracle Linux 8.3
0 GPUs

```
%%time

# alignment functions
from bigstream.align import alignment_pipeline
from bigstream.transform import apply_transform

# define alignment steps
common_kwargs = {
    'alignment_spacing':1.0,
    'shrink_factors':(4,2,1,),
    'smooth_sigmas':(4.,2.,1.,),
    'optimizer_args':{
        'learningRate':0.25,
        'minStep':0.,
        'numberOfIterations':400,
    },
}

affine_kwargs = {
    'alignment_spacing':0.5,
    'shrink_factors':(8,4,2,1),
    'smooth_sigmas':(4.,4.,2.,1.),
}

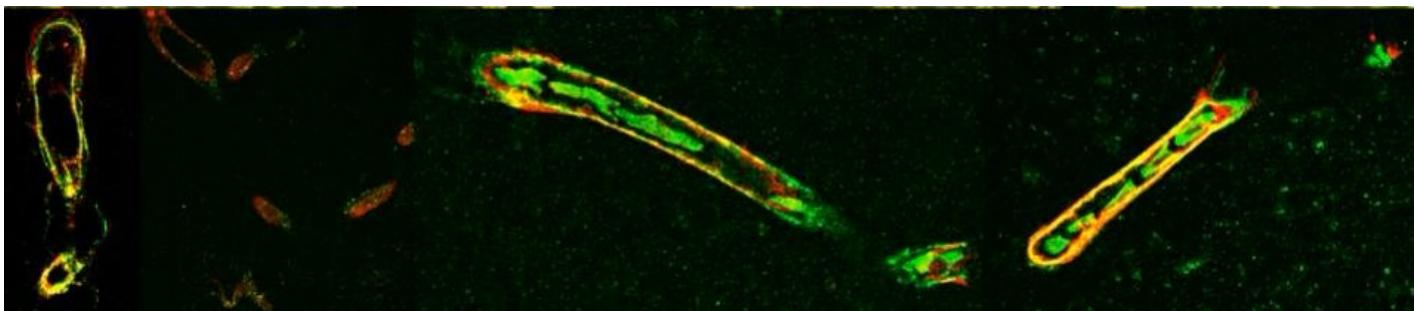
deform_kwargs = {
    'control_point_spacing':32.0,
    'control_point_levels':(1,2,2),
    'fix_mask':fix_mask,
    'mov_mask':mov_mask,
    'optimizer_args':{
        'learningRate':10.0,
        'minStep':0.,
        'numberOfIterations':600,
    },
}

steps = [('affine', {**common_kwargs, **affine_kwargs}),
         ('deform', {**common_kwargs, **deform_kwargs})]

# align
affine, deform = alignment_pipeline(
    fix, mov,
    fix_spacing, mov_spacing,
    steps,
    return_format='compressed',
)

# apply affine only
affine_aligned = apply_transform(
    fix, mov,
    fix_spacing, mov_spacing,
    transform_list=[affine],
)

# apply both
deform_aligned = apply_transform(
    fix, mov,
    fix_spacing, mov_spacing,
    transform_list=[affine, deform],
    extrapolate_with_nn=True,
)
```



Training Data: *C. elegans*

RnR-ExM Training Data, Aligned
FIXED images in RED
MOVING images in GREEN
YELLOW indicates overlap

16 cores: 2.7 GHz Intel SkyLake Platinum 8168
Oracle Linux 8.3
0 GPUs

```
%%time

# alignment functions
from bigstream.align import alignment_pipeline
from bigstream.transform import apply_transform

# define alignment steps
common_kwargs = {
    'alignment_spacing':1.0,
    'shrink_factors':(2,),
    'smooth_sigmas':(1.5,),
    'optimizer_args':{
        'learningRate':0.25,
        'minStep':0.,
        'numberOfIterations':400,
    },
}

affine_kwargs = {
    'initial_condition':'CENTER',
}

affine2_kwargs = {
    'fix_mask':fix_mask,
    'mov_mask':mov_mask,
    'optimizer_args':{
        'learningRate':1.0,
        'minStep':0.,
        'numberOfIterations':3600,
    },
}

deform_kwargs = {
    'fix_mask':fix_mask,
    'mov_mask':mov_mask,
    'control_point_spacing':32.0,
    'control_point_levels':(1,),
    'alignment_spacing':1.0,
    'shrink_factors':(2,),
    'smooth_sigmas':(1.5,),
    'optimizer_args':{
        'learningRate':10.0,
        'minStep':0.,
        'numberOfIterations':60,
    },
}

steps = [('affine', {**common_kwargs, **affine_kwargs},),
         ('affine', {**common_kwargs, **affine2_kwargs},),
         ('deform', {**common_kwargs, **deform_kwargs},),]

# align
affine, deform = alignment_pipeline(
    fix, mov,
    fix_spacing, mov_spacing,
    steps,
    return_format='compressed',
)

# apply affine only
affine_aligned = apply_transform(
    fix, mov,
    fix_spacing, mov_spacing,
    transform_list=[affine],
)

# deform aligned
deform_aligned = apply_transform(
    fix, mov,
    fix_spacing, mov_spacing,
    transform_list=[affine, deform],
)
```

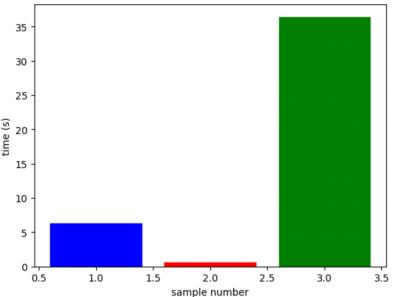
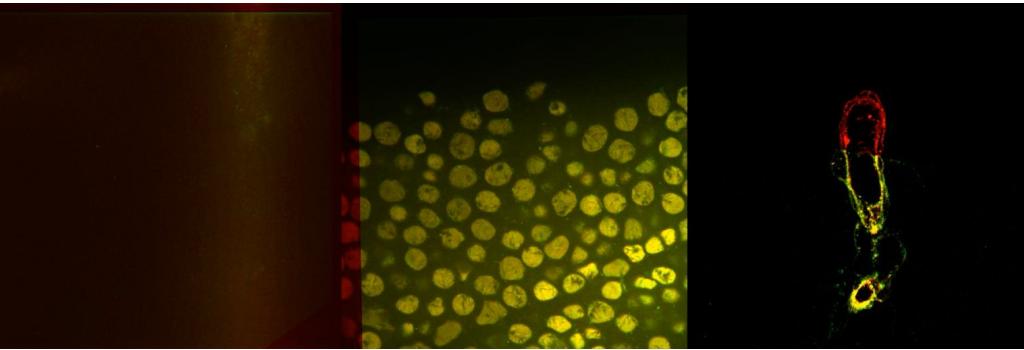
compute time

sample number	time (s)
1	~1750
2	0
3	0
4	0

Lost the compute times for samples 1, 3, and 4

Conclusions from training data

- Bigstream can successfully align all the data types
- Bigstream can handle simple and complex alignments
- Bigstream can align data pairs with significantly different intensity characteristics and SNR
- Multiple alignment strategies are possible with different cost/benefit ratios
- In-memory alignments run efficiently on typical workstation size resources



Validation Data

C elegans

```
# alignment functions
from bigstream.align import alignment_pipeline
from bigstream.transform import apply_transform

# define alignment steps
common_kwargs = {
    'alignment_spacing':2.0,
    'shrink_factors':(2.,),
    'smooth_sigmas':(2.,),
    'optimizer_args':{
        'learningRate':0.25,
        'minStep':0.,
        'numberOfIterations':2000,
    },
}

rigid_kwargs = {
    'initial_condition':'CENTER',
}

steps = [('rigid', {**common_kwargs, **rigid_kwargs}),]

%%time

# run alignment
affine = alignment_pipeline(
    fix, mov,
    fix_spacing, mov_spacing,
    steps,
    # fix_mask=fix_mask,
    # mov_mask=mov_mask,
)
```

Zebrafish

```
# define alignment
from bigstream.align import alignment_pipeline
from bigstream.transform import apply_transform

# define alignment steps
common_kwargs = {
    'alignment_spacing':4.0,
    'shrink_factors':(2.,),
    'smooth_sigmas':(2.,),
    'optimizer_args':{
        'learningRate':0.25,
        'minStep':0.,
        'numberOfIterations':400,
    },
}

affine_kwargs = {
    'initial_condition':'CENTER',
}

steps = [('affine', {**common_kwargs, **affine_kwargs}),]

%%time

# run it
affine = alignment_pipeline(
    fix, mov,
    fix_spacing, mov_spacing,
    steps,
)
```

Mouse

```
# alignment functions
from bigstream.align import alignment_pipeline
from bigstream.transform import apply_transform

# define alignment steps
common_kwargs = {
    'alignment_spacing':0.5,
    'shrink_factors':(8,4,2,1),
    'smooth_sigmas':(4.,4.,2.,1.),
    'optimizer_args':{
        'learningRate':0.25,
        'minStep':0.,
        'numberOfIterations':400,
    },
}

steps = [('affine', common_kwargs),]

%%time

# align
affine = alignment_pipeline(
    fix, mov,
    fix_spacing, mov_spacing,
    steps,
)
```

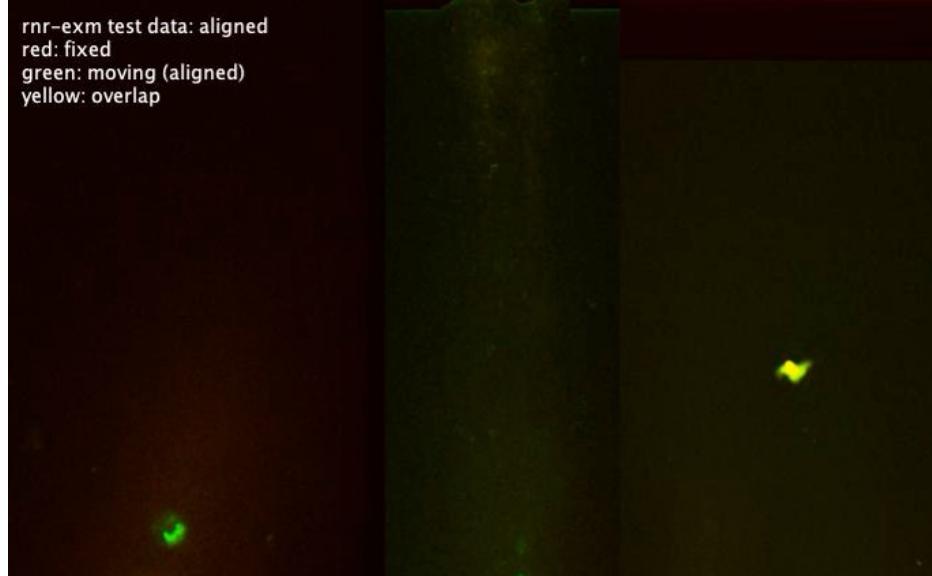
Conclusions from validation data

- Very simple and fast alignment can solve most of the discrepancy for these datasets

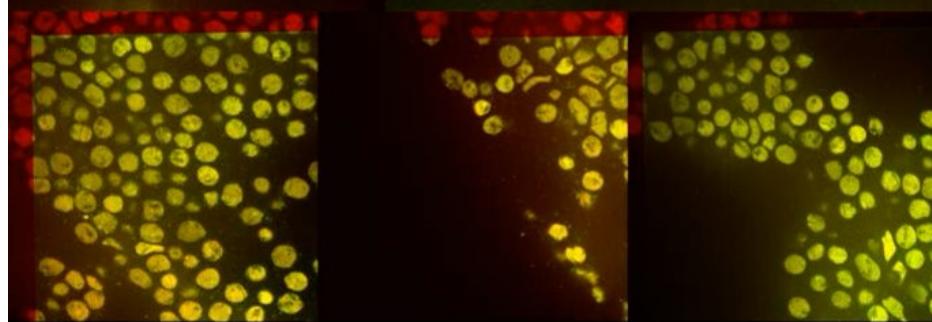
1

Test Data: aligned

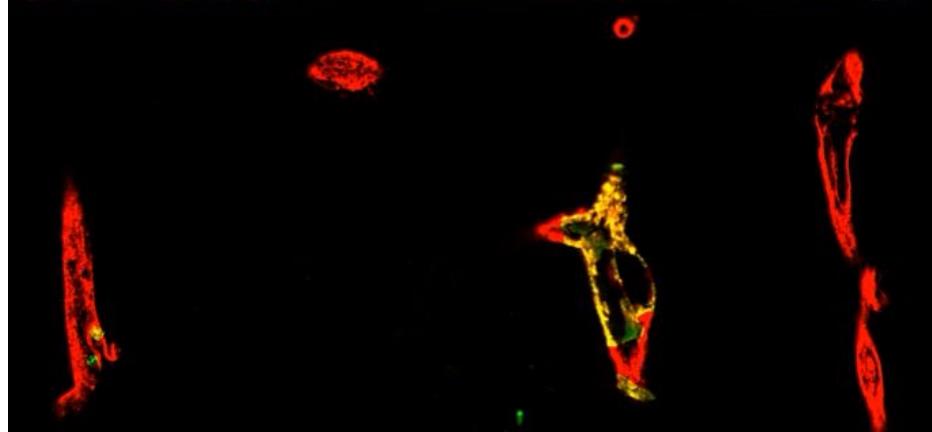
C. elegans



Zebrafish



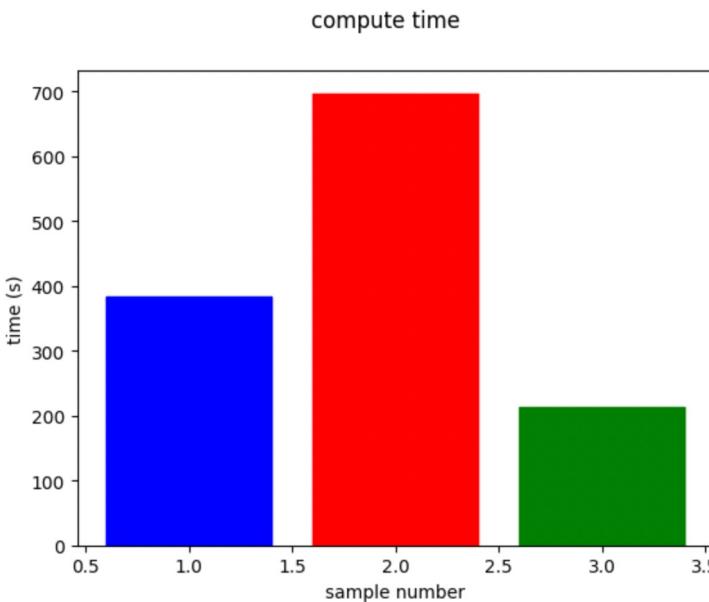
Mouse



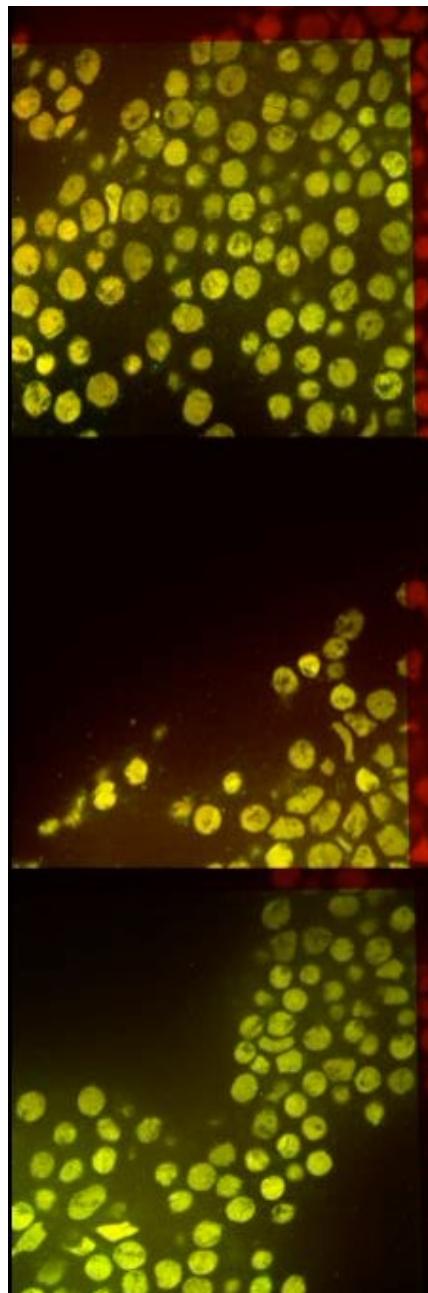
rnr-exm test data: aligned
red: fixed
green: moving (aligned)
yellow: overlap



Test Data: Zebrafish



16 cores: 2.7 GHz Intel SkyLake Platinum 8168
Oracle Linux 8.3
0 GPUs



```
# define alignment
from bigstream.align import alignment_pipeline
from bigstream.transform import apply_transform

# define alignment steps
common_kwargs = {
    'alignment_spacing':0.4,
    'shrink_factors':(2,2,),
    'smooth_sigmas':(2.,0.2,),
    'optimizer_args':{
        'learningRate':0.25,
        'minStep':0.,
        'numberOfIterations':500,
    },
}

affine_kwargs = {
    'initial_condition':'CENTER',
}

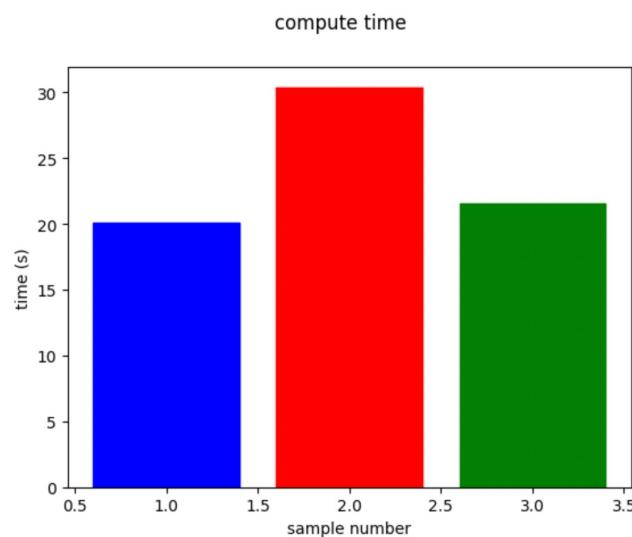
steps = [(['affine', {**common_kwargs, **affine_kwargs}],)

%%time

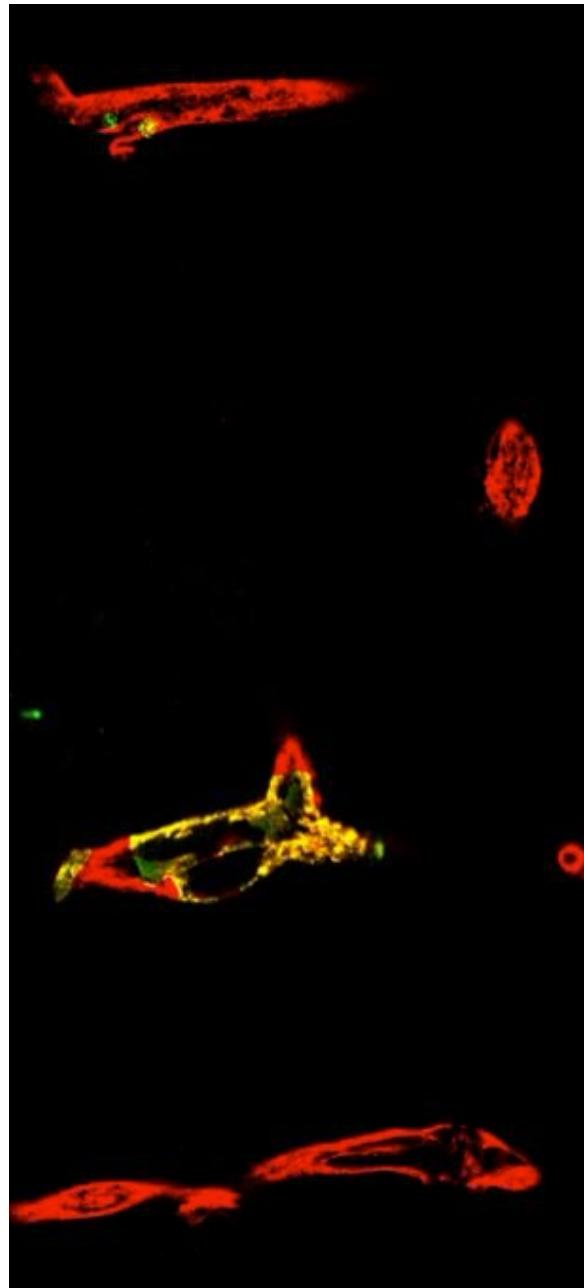
# run it
affine = alignment_pipeline(
    fix, mov,
    fix_spacing, mov_spacing,
    steps,
)
```

BigStream

Test Data: Mouse



16 cores: 2.7 GHz Intel SkyLake Platinum 8168
Oracle Linux 8.3
0 GPUs



```
# alignment functions
from bigstream.align import alignment_pipeline
from bigstream.transform import apply_transform

# define alignment steps
common_kwargs = {
    'alignment_spacing':0.5,
    'shrink_factors':(8,4,2,1),
    'smooth_sigmas':(4.,4.,2.,1.),
    'optimizer_args':{
        'learningRate':0.25,
        'minStep':0.,
        'numberOfIterations':400,
    },
}

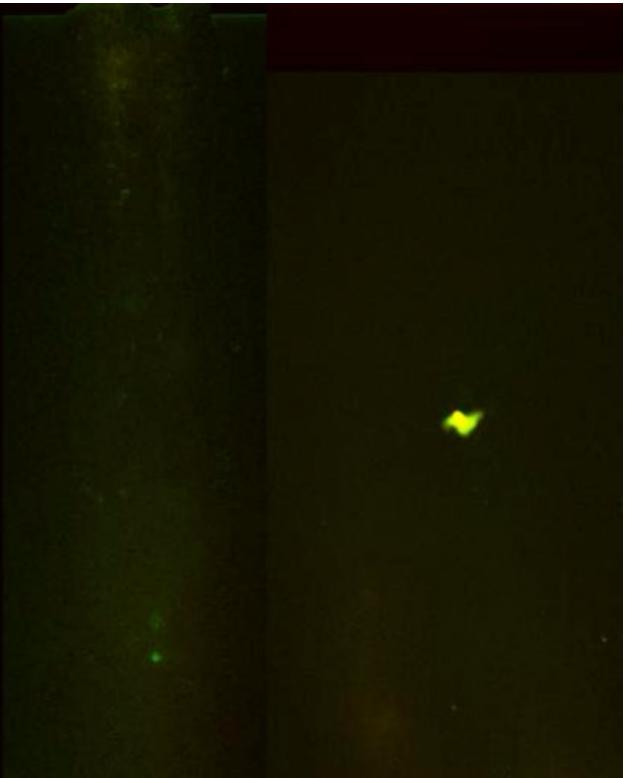
steps = [(['affine', common_kwargs],]

%%time

# align
affine = alignment_pipeline(
    fix, mov,
    fix_spacing, mov_spacing,
    steps,
)
```

Test Data: *C. elegans*

rnr-exm test data: aligned
red: fixed
green: moving (aligned)
yellow: overlap



```
# alignment functions
from bigstream.align import alignment_pipeline
from bigstream.transform import apply_transform
from bigstream.piecewise_align import distributed_piecewise_alignment_pipeline

# define alignment steps
common_kwargs = {
    'alignment_spacing':2.0,
    'shrink_factors':(2,1),
    'smooth_sigmas':(12.,2.,),
    'optimizer_args':{
        'learningRate':0.25,
        'minStep':0.,
        'numberOfIterations':6000,
    },
}

rigid_kwargs = {
    'initial_condition':'CENTER',
}

affine_kwargs = {
    'smooth_sigmas':(2.,2.,),
}

steps = [('rigid', {**common_kwargs, **rigid_kwargs},),
          ('affine', {**common_kwargs, **affine_kwargs},,)]

dpap_affine_kwargs = {
    'alignment_spacing':None,
    'shrink_factors':(1,),
    'smooth_sigmas':(0.4,),
    'optimizer_args':{
        'learningRate':0.25,
        'minStep':0.,
        'numberOfIterations':800,
    },
}

dpap_deform_kwargs = {
    'control_point_spacing':3.0,
    'control_point_levels':(1,),
    'alignment_spacing':None,
    'shrink_factors':(1,),
    'smooth_sigmas':(0.1625,),
    'optimizer_args':{
        'learningRate':0.25,
        'minStep':0.,
        'numberOfIterations':200,
    },
}

dpap_steps = [(['affine', dpap_affine_kwargs,,),
              ('deform', dpap_deform_kwargs,,)]
```

```
cluster_kwargs = {
    'project':'ahrens',
    'ncpus':1,
    'threads':1,
    'min_workers':200,
    'max_workers':800,
}

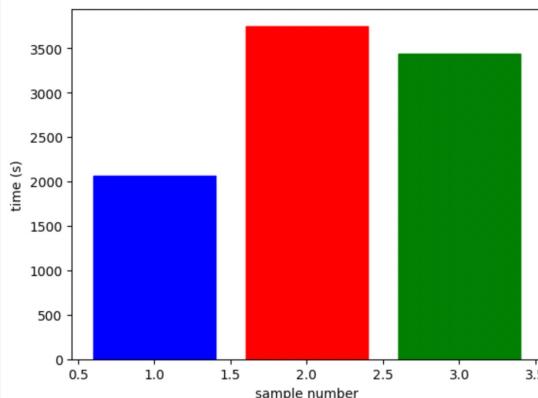
blocksize = [24, 48, 48]

%%time

# run alignment
affine = alignment_pipeline(
    fix, mov,
    fix_spacing, mov_spacing,
    steps,
    fix_mask=fix_mask,
    mov_mask=mov_mask,
)

# run piecewise alignment
deform = distributed_piecewise_alignment_pipeline(
    fix, mov,
    fix_spacing, mov_spacing,
    dpap_steps,
    blocksize,
    fix_mask=fix_mask,
    mov_mask=mov_mask,
    metric='C',
    static_transform_list=[affine,],
    cluster_kwargs=cluster_kwargs,
)
```

compute time



800 cores: 2.7 GHz Intel SkyLake Platinum 8168
Oracle Linux 8.3
0 GPUs

Conclusions from test data

- Simple and quick alignments are still very effective
- Costly distributed alignments are possible and can give increased precision if you want to pay the cost

Conclusions from the overall challenge

- DSC is a good metric, but we all overfit to the single annotation
- Both traditional methods and learning methods can achieve best performance but with different cost/benefit trade offs
 - This is very fruitful discussion topic in person and in the literature
- Challenges are fun and push the sharing of tools and advance the state of the art

Give Bigstream a try! github.com/GFleishman/bigstream

I've shown some of the general features of the package but there are a lot more available